

# Introduction to Data Engineering

## 07 Clustering, Outlier, Novelty Detection and Duplicate Detection

Junbo Huang and Angelie Kraft and Cedric Möller and David Rath and Ricardo Usbeck

<https://lernen.min.uni-hamburg.de/course/view.php?id=2917>

# Intuition: Aspect of Human Learning

- How do we learn? -> how do children learn?

# Intuition: Aspect of Human Learning

- How do we learn? -> how do children learn?
- Behaviorists: conditioning -> Reinforcement learning

# Intuition: Aspect of Human Learning

- How do we learn? -> how do children learn?
- Behaviorists: conditioning -> Reinforcement learning
- Albert Bandura: social learning theory -> Imitation learning

# Intuition: Aspect of Human Learning

- How do we learn? -> how do children learn?
- Behaviorists: conditioning -> Reinforcement learning
- Albert Bandura: social learning theory -> Imitation learning
- Lower-level learning?
  - ▶ Perception? Cognition? Concept learning?

# Intuition: Aspect of Human Learning

- How do we learn? -> how do children learn?
- Behaviorists: conditioning -> Reinforcement learning
- Albert Bandura: social learning theory -> Imitation learning
- Lower-level learning?
  - ▶ Perception? Cognition? Concept learning?
  - ▶ Wittgenstein, 1953: language acquisition through pointing<sup>1</sup>

---

<sup>1</sup>Wittgenstein, L. (1953). Philosophical investigations.

# Intuition: Aspect of Human Learning

- How do we learn? -> how do children learn?
- Behaviorists: conditioning -> Reinforcement learning
- Albert Bandura: social learning theory -> Imitation learning
- Lower-level learning?
  - ▶ Perception? Cognition? Concept learning?
  - ▶ Wittgenstein, 1953: language acquisition through **pointing**<sup>2</sup>.
  - ▶ Eleanor Rosch, 1978: Category formation is strongly connected to forming **prototypical concepts**<sup>3</sup>.

---

<sup>2</sup>Wittgenstein, L. (1953). Philosophical investigations.

<sup>3</sup>Rosch, E. (1978). Principles of categorization.

# Intuition: Aspect of Human Learning

- The concept of birdiness.
- Robin is a more typical bird than a penguin.




Figure 1 Birdiness rankings

Source: Aitchison, J. (1994). Words in the mind: An introduction to the mental lexicon. page 54.



# Intuition: Aspect of Human Learning

- The concept of birdiness.
- Robin is a more typical bird than a penguin.
- Bias in prototyping?
  - ▶ Misunderstanding.




Figure 1 Birdiness rankings


Source: Aitchison, J. (1994). Words in the mind: An introduction to the mental lexicon. page 54.

# Overview

- 1 Distance Measures
  - Distance Between Points
  - Distance Between Clusters
- 2 Clustering
- 3 Clustering
  - Hard Clustering
  - Soft Clustering
- 4 Outliers, Novelty and Duplicate Detection


# Two Types of Distance Measures

What are the two types of distance measures?



# Two Types of Distance Measures

- Distance between points
- Distance between clusters



# Distance Metric

Consider a metric space  $\mathcal{X}$ . A **distance metric** is a mapping  $d : \mathcal{X} \times \mathcal{X} \rightarrow [0, \infty)$  which satisfies the following properties:

- non-negativity:  $d(x_i, x_j) \geq 0$
- identity:  $d(x_i, x_j) = 0 \Leftrightarrow x_i = x_j$
- symmetry:  $d(x_i, x_j) = d(x_j, x_i)$
- triangle inequality:  $d(x_i, x_j) \leq d(x_i, x_k) + d(x_k, x_j)$

where  $x_i, x_j, x_k \in \mathcal{X}$ .

# Distance Metric

## Remarks:

- In Mathematics, the term **metric** is used only when the axioms are fulfilled.
- In ML, **distance metric** often refers to the similarity or dissimilarity measure and it may not satisfy all the axioms (e.g., cosine distance).
- Therefore, if you are clear that the axioms are satisfied, use the term **metric**.

# Distance Between Points

Given two points  $x, y \in \mathbb{R}^n$ ,


- Euclidean Distance:

$$d(x, y) = \sqrt{\sum_i^n (x_i - y_i)^2}$$

- Manhattan Distance / City Block Distance:

$$d(x, y) = \sum_i^n |x_i - y_i|$$

- (Any parametrized distance metric)



# Distance Between Clusters


How can one measure distance between two clusters?



# Distance Between Clusters


Given two clusters  $X$  and  $Y$ :

- $D_{min}(X, Y) = \min_{x \in X, y \in Y} d(x, y)$
- $D_{max}(X, Y) = \max_{x \in X, y \in Y} d(x, y)$
- $D_{mean}(X, Y) = \frac{1}{|X||Y|} \sum_{x \in X, y \in Y} d(x, y)$
- $D_{centroid}(X, Y) = d(\frac{1}{|X|} \sum_{x \in X} x, \frac{1}{|Y|} \sum_{y \in Y} y)$




# Distance Between Clusters

$D_{min}$ ,  $D_{max}$ ,  $D_{mean}$ ,  $D_{centroid}$  ignore the variances of data.



# Distance Between Clusters

$D_{min}$ ,  $D_{max}$ ,  $D_{mean}$ ,  $D_{centroid}$  ignore the variances of data.



- Solution: Mahalanobis Distance

# Mahalanobis Distance

Idea: Scaling of distances using the covariance matrix  $\Sigma \in \mathbb{R}^{d \times d}$ .

$$D_{\Sigma}(X, Y) = \sqrt{(x - y)^T \Sigma^{-1} (x - y)}$$

Properties:

- $\Sigma$  characterizes the distribution of data.
- If  $\Sigma$  is an identity matrix, the above equation represents Euclidean distance.
- It is a good distance measure between sets/clusters.
- Can also be used to measure distances between a point and sets/clusters.
- Spectral distance metric learning: parametrizing  $\Sigma$ .

# A Clustering Problem

Clustering German cities based on the location using Euclidean distance.

- Hamburg
- Bremen
- Hannover
- Leipzig
- Frankfurt
- Nürnberg



Source: Google maps

# A Clustering Problem

$$\mathbf{x} = \begin{bmatrix} 53.55 & 9.99 \\ 53.07 & 8.8 \\ 52.37 & 9.73 \\ 51.33 & 12.37 \\ 50.13 & 8.66 \\ 49.45 & 11.07 \end{bmatrix}$$

Or, a  $n \times d$  matrix ( $n$  cities and  $d$  features)

# Distance Matrix

A matrix where each entry represents the euclidean distance (in km) between two cities.

	Hamburg	Bremen	Hannover	Leipzig	Frankfurt	Nürnberg
Hamburg	0	95	133	294	393	462
Bremen		0	100	310	330	433
Hannover			0	214	262	338
Leipzig				0	293	229
Frankfurt					0	187
Nürnberg						0

# Types of Clustering Algorithm

- Hard clustering algorithms
  - ▶ Hard assignment of clusters
  - ▶ Hierarchical clustering, k-means<sup>4</sup>...
- Soft clustering algorithms
  - ▶ Probabilistic assignment of clusters
  - ▶ Gaussian Mixture Models...

---

<https://colab.research.google.com/drive/1x95aarcKwFMs26u4Fwn5rQThWINHhQ65?usp=sharing>



# Hierarchical Clustering

- Do not need to pre-define number of clusters  $k$ .
- Two complementary methods:
  - ▶ Agglomerative Clustering: bottom-up
  - ▶ Divisive Clustering: top-down

# Agglomerative Clustering

Basic agglomerative clustering:

- ➊ Assign each object to its own single-object cluster.
- ➋ Choose the closest pair of clusters and merge them into a single cluster.
- ➌ Calculate the distance between the new cluster and each of the old clusters.
- ➍ Assign Repeat steps 2 and 3 until all the objects are in a single cluster.

# Agglomerative Clustering

	Hamburg	Bremen	Hannover	Leipzig	Frankfurt	Nürnberg
Hamburg	0	95	133	294	393	462
Bremen		0	100	310	330	433
Hannover			0	214	262	338
Leipzig				0	293	229
Frankfurt					0	187
Nürnberg						0

① {Hamburg}, {Bremen}, {Hannover}, {Leipzig}, {Frankfurt}, {Nürnberg}

# Agglomerative Clustering

	Hamburg	Bremen	Hannover	Leipzig	Frankfurt	Nürnberg
Hamburg	0	95	133	294	393	462
Bremen		0	100	310	330	433
Hannover			0	214	262	338
Leipzig				0	293	229
Frankfurt					0	187
Nürnberg						0

- 1 {Hamburg}, {Bremen}, {Hannover}, {Leipzig}, {Frankfurt}, {Nürnberg}
- 2 {Hamburg, Bremen}, {Hannover}, {Leipzig}, {Frankfurt}, {Nürnberg}

# Agglomerative Clustering

	Hamburg	Bremen	Hannover	Leipzig	Frankfurt	Nürnberg
Hamburg	0	95	133	294	393	462
Bremen		0	100	310	330	433
Hannover			0	214	262	338
Leipzig				0	293	229
Frankfurt					0	187
Nürnberg						0

- 1 {Hamburg}, {Bremen}, {Hannover}, {Leipzig}, {Frankfurt}, {Nürnberg}
- 2 {Hamburg, Bremen}, {Hannover}, {Leipzig}, {Frankfurt}, {Nürnberg}
- 3 {Hamburg, Bremen, Hannover}, {Leipzig}, {Frankfurt}, {Nürnberg}

# Agglomerative Clustering

	Hamburg	Bremen	Hannover	Leipzig	Frankfurt	Nürnberg
Hamburg	0	95	133	294	393	462
Bremen		0	100	310	330	433
Hannover			0	214	262	338
Leipzig				0	293	229
Frankfurt					0	187
Nürnberg						0

- 1 {Hamburg}, {Bremen}, {Hannover}, {Leipzig}, {Frankfurt}, {Nürnberg}
- 2 {Hamburg, Bremen}, {Hannover}, {Leipzig}, {Frankfurt}, {Nürnberg}
- 3 {Hamburg, Bremen, Hannover}, {Leipzig}, {Frankfurt}, {Nürnberg}
- 4 {Hamburg, Bremen, Hannover}, {Leipzig}, {Frankfurt, Nürnberg}

# Agglomerative Clustering

	Hamburg	Bremen	Hannover	Leipzig	Frankfurt	Nürnberg
Hamburg	0	95	133	294	393	462
Bremen		0	100	310	330	433
Hannover			0	214	262	338
Leipzig				0	293	229
Frankfurt					0	187
Nürnberg						0

- 1 {Hamburg}, {Bremen}, {Hannover}, {Leipzig}, {Frankfurt}, {Nürnberg}
- 2 {Hamburg, Bremen}, {Hannover}, {Leipzig}, {Frankfurt}, {Nürnberg}
- 3 {Hamburg, Bremen, Hannover}, {Leipzig}, {Frankfurt}, {Nürnberg}
- 4 {Hamburg, Bremen, Hannover}, {Leipzig}, {Frankfurt, Nürnberg}
- 5 {Hamburg, Bremen, Hannover, Leipzig}, {Frankfurt, Nürnberg}

# Agglomerative Clustering

	Hamburg	Bremen	Hannover	Leipzig	Frankfurt	Nürnberg
Hamburg	0	95	133	294	393	462
Bremen		0	100	310	330	433
Hannover			0	214	262	338
Leipzig				0	293	229
Frankfurt					0	187
Nürnberg						0


- 1 {Hamburg}, {Bremen}, {Hannover}, {Leipzig}, {Frankfurt}, {Nürnberg}
- 2 {Hamburg, Bremen}, {Hannover}, {Leipzig}, {Frankfurt}, {Nürnberg}
- 3 {Hamburg, Bremen, Hannover}, {Leipzig}, {Frankfurt}, {Nürnberg}
- 4 {Hamburg, Bremen, Hannover}, {Leipzig}, {Frankfurt, Nürnberg}
- 5 {Hamburg, Bremen, Hannover, Leipzig}, {Frankfurt, Nürnberg} -> Single-linkage clustering
- 5 {Hamburg, Bremen, Hannover}, {Leipzig, Frankfurt, Nürnberg} -> Complete-linkage clustering



## Recall: Distance Between Clusters

Given two clusters  $X$  and  $Y$ :

- $D_{min}(X, Y) = \min_{x \in X, y \in Y} d(x, y)$ 
  - ▶ single-linkage clustering
- $D_{max}(X, Y) = \max_{x \in X, y \in Y} d(x, y)$ 
  - ▶ complete-linkage clustering
- $D_{mean}(X, Y) = \frac{1}{|X||Y|} \sum_{x \in X, y \in Y} d(x, y)$
- $D_{centroid}(X, Y) = d(\frac{1}{|X|} \sum_{x \in X} x, \frac{1}{|Y|} \sum_{y \in Y} y)$




# Agglomerative Clustering

	Hamburg	Bremen	Hannover	Leipzig	Frankfurt	Nürnberg
Hamburg	0	95	133	294	393	462
Bremen		0	100	310	330	433
Hannover			0	214	262	338
Leipzig				0	293	229
Frankfurt					0	187
Nürnberg						0

- 1 {Hamburg}, {Bremen}, {Hannover}, {Leipzig}, {Frankfurt}, {Nürnberg}
- 2 {Hamburg, Bremen}, {Hannover}, {Leipzig}, {Frankfurt}, {Nürnberg}
- 3 {Hamburg, Bremen, Hannover}, {Leipzig}, {Frankfurt}, {Nürnberg}
- 4 {Hamburg, Bremen, Hannover}, {Leipzig}, {Frankfurt, Nürnberg}
- 5 {Hamburg, Bremen, Hannover, Leipzig}, {Frankfurt, Nürnberg}
- 6 {Hamburg, Bremen, Hannover, Leipzig, Frankfurt, Nürnberg}


# Agglomerative Clustering

Dendrogram: a diagram representing a binary tree.




# Agglomerative Clustering

Dendrogram: a diagram representing a binary tree.




# Agglomerative Clustering

Dendrogram: a diagram representing a binary tree.



# Agglomerative Clustering

Dendrogram: a diagram representing a binary tree.



# Agglomerative Clustering

Properties of agglomerative clustering?

# Agglomerative Clustering


## Properties:

- Any distance measure can be used
- we only need the distance matrix
- No parameters
- Resulting dendrogram needs to be analyzed to decide number of desired clusters
- Slow when the number of samples is large
- Greedy, deterministic



# K-means Algorithm

- 1 Randomly choose  $k$  examples as initial centroids.
- 2 Create  $k$  clusters by assigning each example to its closest centroid.
- 3 Compute  $k$  new centroids by averaging examples in each cluster.
- 4 Repeat step 2 and 3 until centroids don't change.




# How to choose k?

Possible solutions:

- Run k-means multiple times with random initialization.
- Take a subset of the data, and run hierarchical clustering.
- Statistics
  - ▶ The elbow method
    - ★ Calculate the Within-Cluster-Sum of Squared Errors (WCSS) for different values of k
  - ▶ The Silhouette Method
    - ★ Measure how similar a point is to its own cluster (cohesion) compared to other clusters (separation)

# Properties of K-means

- Only one parameter  $k$ .
  - ▶ Implicitly defines scale and resulting shape of clusters.
- Fast.
- Greedy, non-deterministic -> local optima, depending on initial conditions.




# Problems of K-means

- Sensitive to outliers.
  - ▶ Outliers highly influence clustering result in naïve k-means;
  - ▶ Outliers can be handled in the preprocessing step, or with algorithms which are robust to outliers, e.g., DBSCAN.
- What if we don't know  $k$ ?
  - ▶ Other algorithms, e.g., hierarchical clustering or DBSCAN.
- What if data does not depict circular shape?
  - ▶ If you know how the data should look like,
  - ▶ You can choose other algorithms that have a more relaxed constraint on the shape of clusters, e.g., soft clustering or DBSCAN.

# DBSCAN Clustering


- Ester, 1996
- No need to define  $k$
- Density-based clustering algorithm
  - ▶ Density connected points belong to the same cluster
- Not entirely deterministic.
- Two parameters:
  - ▶ epsilon  $\epsilon$
  - ▶ Minimum number of points **MinPts**



# DBSCAN Clustering

Steps:


- 1 Find all core points given  $\epsilon$  and *MinPts*.
- 2 For each core point if it is not already assigned to a cluster, create a new cluster.
- 3 Find recursively all its density connected points and assign them to the same cluster as the core point.
- 4 Assign points that do not belong to any cluster as noise/ outliers.



# DBSCAN Clustering

Steps:


- 1 Find all core points given  $\epsilon$  and *MinPts*.
- 2 For each core point if it is not already assigned to a cluster, create a new cluster.
- 3 Find recursively all its density connected points and assign them to the same cluster as the core point.
- 4 Assign points that do not belong to any cluster as noise/ outliers.



# DBSCAN Clustering

Steps:

- 1 Find all core points given  $\epsilon$  and *MinPts*.
- 2 For each core point if it is not already assigned to a cluster, create a new cluster.
- 3 Find recursively all its density connected points and assign them to the same cluster as the core point.
- 4 Assign points that do not belong to any cluster as noise/ outliers.






# DBSCAN Clustering

Steps:


- 1 Find all core points given  $\epsilon$  and *MinPts*.
- 2 For each core point if it is not already assigned to a cluster, create a new cluster.
- 3 Find recursively all its density connected points and assign them to the same cluster as the core point.
- 4 Assign points that do not belong to any cluster as noise/ outliers.



# DBSCAN Clustering

Steps:

- 1 Find all core points given  $\epsilon$  and *MinPts*.
- 2 For each core point if it is not already assigned to a cluster, create a new cluster.
- 3 Find recursively all its density connected points and assign them to the same cluster as the core point.
- 4 Assign points that do not belong to any cluster as noise/ outliers.




# DBSCAN Clustering

Pros:

- 1 Robust to outliers.
- 2 No need to define  $k$
- 3 Arbitrary cluster shape.

Cons:

- 1 Difficult to choose  $\epsilon$  if data is not well understood.
- 2 "curse of dimensionality".
  - 1 a good distance measure is difficult.




# DBSCAN Clustering

Problems of DBSCAN?

# DBSCAN Clustering

One problem:

- How to cluster a point between two clusters?
- Both k-means and DBSCAN assign one cluster to each data point.
- There is no "likelihood" in cluster assignments.



# Soft Clustering

So far,

- clusters are described by a set of data points or by the centers  
-> **clusters are disjoint**
- This is called hard clustering.
- Each data point is assigned to a single cluster.
- No way to express uncertainty/confidence about the assignment to a cluster.


Now,

- Describe data by probability distribution  $P(x)$ .
- **Soft clustering:**
- A data point is assigned to the clusters with uncertainty/confidence.
- As a result, clusters do not have hard boundaries.

# Soft clustering

Let's think about soft assignment in Deep Learning...

- What is the probability distribution outputted by a neural network?




Source: Feron, B., Moraleda, G., Vossen, J., & Monti, A. Probabilistic Short-Term Residential Load Forecasting based on Feed Forward and LSTM Recurrent Neural Networks.

# Soft clustering

Let's think about soft assignment in Deep Learning...

- What is the probability distribution outputted by a neural network?
- $\text{Softmax}(x_j) = \frac{\exp(x_j)}{\sum_j \exp(x_j)}$
- What is probability?




Source: Feron, B., Moraleda, G., Vossen, J., & Monti, A. Probabilistic Short-Term Residential Load Forecasting based on Feed Forward and LSTM Recurrent Neural Networks.



# Soft clustering

Let's think about soft assignment in Deep Learning...

- What is the probability distribution outputted by a neural network?
- $\text{Softmax}(x_i) = \frac{\exp(x_i)}{\sum_j \exp(x_j)}$
- Probability  $\neq$  Softmax output<sup>5</sup>



Source: Feron, B., Moraleda, G., Vossen, J., & Monti, A. Probabilistic Short-Term Residential Load Forecasting based on Feed Forward and LSTM Recurrent Neural Networks.

<sup>5</sup>Guo, C., Pleiss, G., Sun, Y., & Weinberger, K. Q. (2017, July). On calibration of modern neural networks. In International Conference on Machine Learning (pp. 1321-1330). PMLR.


# Soft clustering

## Gaussian Mixture Models (GMMs)

- Assign a Gaussian to each cluster center.
- Linear superposition of  $K$  Gaussians.

$$p(x|\theta) = \sum_{k=1}^K \pi_k \mathcal{N}(x|\mu_k, \Sigma_k)$$

- $\mathcal{N}(\cdot)$  denotes a Gaussian with mean  $\mu$  and covariate matrix  $\Sigma$ .
- $\pi_k$  denotes the *a priori* probability that a data point belongs to cluster  $k$ .



# Gaussian Mixture Models

A density model where we combine a finite number of  $K$  Gaussian distributions  $\mathcal{N} = (\mathbf{x}|\mu_k, \Sigma_k)$  so that

$$p(\mathbf{x}|\theta) = \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x}|\mu_k, \Sigma_k)$$

$$0 \leq \pi_k \leq 1, \sum_{k=1}^K \pi_k = 1$$

- $\theta := \{\pi_k, \mu_k, \Sigma_k : k = 1, 2, \dots, K\}$
- Goal: optimize  $\theta$  such that the GMM fits the data (Maximum Likelihood Estimate).

# Gaussian Mixture Models

## Problem formulation:

- Given dataset  $\mathcal{X} = \{x_1, x_2, \dots, x_N\}$ ,  $n = 1, \dots, N$ , where  $x_n$  are drawn i.i.d. from an unknown distribution  $p(x)$ . Our objective is to find a good approximation of this unknown distribution  $p(x)$  by means of a GMM with  $K$  mixture components, given a set of parameters  $\theta := \{\pi_k, \mu_k, \Sigma_k : k = 1, 2, \dots, K\}$ .
- Typically done by likelihood maximization (or negative log likelihood minimization).

$$p(\mathcal{X}|\theta) = \prod_{n=1}^N p(x_n|\theta), \text{ where } p(x_n|\theta) = \sum_{k=1}^K \pi_k \mathcal{N}(x_n|\mu_k, \Sigma_k)$$

$$\mathcal{L} = -\log p(\mathcal{X}|\theta) = -\sum_{n=1}^N \log p(x_n|\theta) = -\sum_{n=1}^N \log \sum_{k=1}^K \pi_k \mathcal{N}(x_n|\mu_k, \Sigma_k)$$

- Goal: find  $\theta^*$  that minimizes the loss (negative log-likelihood)  $\mathcal{L}$ .

# Gaussian Mixture Models

An important concept: *Responsibility*

$$r_{nk} := \frac{\pi_k \mathcal{N}(x_n | \mu_k, \Sigma_k)}{\sum_{j=1}^K \pi_j \mathcal{N}(x_n | \mu_j, \Sigma_j)}$$

- is defined as the *responsibility* of the  $k^{th}$  mixture component for data point  $x_n$ .
- *Responsibility* is the soft assignment of clusters.
- To update  $r_{nk}$ , we need  $\{\pi_k, \mu_k, \Sigma_k : k = 1, 2, \dots, K\}$ .
- To update  $\{\pi_k, \mu_k, \Sigma_k : k = 1, 2, \dots, K\}$ , we need  $r_{nk}$ .
- Needs an iterative solution.

# Expectation-Maximization Algorithm

- Initialize  $\pi_k, \mu_k, \Sigma_k$ .
- *E-step*: Evaluate responsibilities  $r_{nk}$  for every data point  $x_n$  using current parameters  $\pi_k, \mu_k, \Sigma_k$ :

$$r_{nk} := \frac{\pi_k \mathcal{N}(x_n | \mu_k, \Sigma_k)}{\sum_{j=1}^K \pi_j \mathcal{N}(x_n | \mu_j, \Sigma_j)}$$

- *M-step*: Re-estimate parameters  $\pi_k, \mu_k, \Sigma_k$  using the current responsibilities  $r_{nk}$  (from E-step):

$$\mu_k = \frac{1}{N_k} \sum_{n=1}^N r_{nk} x_n, \quad \pi_k = \frac{N_k}{N},$$
$$\Sigma_k = \frac{1}{N_k} \sum_{n=1}^N r_{nk} (x_n - \mu_k)(x_n - \mu_k)^T.$$

- Repeat E-step and M-step until convergence.

# Expectation-Maximization Algorithm

## Mean updates $\mu_k$

- The mean  $\mu_k$  is pulled toward a data point  $x_n$  with the strength given by  $r_{nk}$ .

$$\mu_k = \frac{1}{N_k} \sum_{n=1}^N r_{nk} x_n$$


- $N_k$  is defined as the total responsibility of the  $k^{th}$  mixture component for the entire dataset.

$$N_k := \sum_{n=1}^N r_{nk}$$


- Importance-weighted Monte Carlo estimate of the mean.

# Expectation-Maximization Algorithm

Mean updates  $\mu_k$



(a) GMM density and individual components prior to updating the mean values.



(b) GMM density and individual components after updating the mean values.

Source: Mathematics for Machine Learning. page 355.



# Expectation-Maximization Algorithm

## Covariance updates $\Sigma_k$


- The covariance matrix is re-estimated based on the new responsibilities  $r_k$

$$\Sigma_k = \frac{1}{N_k} \sum_{n=1}^N r_{nk} (x_n - \mu_k)(x_n - \mu_k)^T$$


- Importance-weighted covariance of data points  $x_n$  associated with the  $k^{th}$  mixture component.
- Proof:
  - ▶ Take the partial derivative of the log-likelihood with respect to  $\Sigma_k$ .
  - ▶ Set the partial derivative to 0.

# Expectation-Maximization Algorithm

Covariance updates  $\Sigma_k$



(a) GMM density and individual components prior to updating the variances.



(b) GMM density and individual components after updating the variances.

Source: Mathematics for Machine Learning. page 358.

# Expectation-Maximization Algorithm

Mixture weights updates  $\pi_k$


- The mixture weights are updated based on the new responsibilities  $r_k$

$$\pi_k = \frac{N_k}{N}$$


- $N$  denotes number of data points.
- The ratio of the total responsibility of the  $k^{th}$  cluster and the number of data points.

# Expectation-Maximization Algorithm

## Mixture weights updates $\pi_k$




(a) GMM density and individual components prior to updating the mixture weights.




(b) GMM density and individual components after updating the mixture weights.

Source: Mathematics for Machine Learning. page 360.


# Expectation-Maximization Algorithm




(c) EM initialization.



(d) EM after one iteration.



(e) EM after 10 iterations.



(f) EM after 62 iterations.


Source: Mathematics for Machine Learning. page 362.

# New stuff?

- DeepCluster (Caron et al. 2018)
- Deep k-means (Fard, Thonet, and Gaussier 2020)
- Swapping Assignments between multiple Views (SwAV) (Caron et al. 2020)
- Prototypical contrastive learning (Li et al. 2020)

# Outliers, Novelty and Duplicate Detection

- Anomaly detection
  - ▶ Outlier detection,
  - ▶ Novelty detection.
- Samples located in low-density region.
- Outliers typically pollutes the data,
- whereas novelties gain insight about the data.
- Methods:
  - ▶ DBSCAN,
  - ▶ Mahalanobis distance, etc.



# Summary

- Different types of distance metrics and their characteristics.
  - ▶ Distance metrics between points,
  - ▶ Distance metrics between clusters.
- Hard clustering algorithms.
  - ▶ Hierarchical clustering, k-means, and DBSCAN.
  - ▶ Easy to use,
  - ▶ But no confidence measurement; can be slow due to "curse of dimensionality".
- Soft clustering algorithms.
  - ▶ Gaussian mixture models, Expectation-Maximization algorithm.
  - ▶ Soft assignment of clusters with uncertainty.
- Application: anomaly detection.

For next time, please read about data quality <https://5stardata.info/en/>



# References

- Ester, M., Kriegel, H. P., Sander, J., Xu, X. (1996, August). A density-based algorithm for discovering clusters in large spatial databases with noise. In kdd (Vol. 96, No. 34, pp. 226-231).
- Deisenroth, M. P., Faisal, A. A., Ong, C. S. (2020). Mathematics for machine learning. Cambridge University Press.
- Ozdemir, S. (2016). Principles of data science. Packt Publishing Ltd.
- Ghojogh, B., Ghodsi, A., Karray, F., Crowley, M. (2022). Spectral, Probabilistic, and Deep Metric Learning: Tutorial and Survey. arXiv preprint arXiv:2201.09267.
- Caron, M., Bojanowski, P., Joulin, A., Douze, M. (2018). Deep clustering for unsupervised learning of visual features. In Proceedings of the European conference on computer vision (ECCV) (pp. 132-149).

# References

- Fard, M. M., Thonet, T., Gaussier, E. (2020). Deep k-means: Jointly clustering with k-means and learning representations. Pattern Recognition Letters, 138, 185-192.
- Caron, M., Misra, I., Mairal, J., Goyal, P., Bojanowski, P., Joulin, A. (2020). Unsupervised learning of visual features by contrasting cluster assignments. Advances in Neural Information Processing Systems, 33, 9912-9924.
- la Rosa, L. E. C., Oliveira, D. A. B. (2022). Learning from Label Proportions with Prototypical Contrastive Clustering.
- Li, J., Zhou, P., Xiong, C., Hoi, S. C. (2020). Prototypical contrastive learning of unsupervised representations. arXiv preprint arXiv:2005.04966.
- Guo, C., Pleiss, G., Sun, Y., Weinberger, K. Q. (2017, July). On calibration of modern neural networks. In International Conference on Machine Learning (pp. 1321-1330). PMLR.